

# Build your Microservices with IBM Kubernetes & Istio



Red Hat



# Agenda

---

- Hands-on preparation 10 min
- Cloud-native development and Istio – presentation 20 min
- Hands-on / Demo – 1:20
- Quiz 10 min
  
- Total time: 2 hours

<https://github.com/agavrin/cloud-native-workshop-2021>



# Check your email for credentials

---

**IBM Workshops** <ibmcloudcoupon.noreply@gmail.com>

кому: я ▼

Hello A!

Please use the following assigned resource:

**server=158.177.16.131**


**userid=user1**

**password=8MPVQyhR63xao**

Any questions - please contact: [agavrin@ru.ibm.com](mailto:agavrin@ru.ibm.com)

# Characteristics of **cloud-native** applications

- Applications that adapt to the Cloud (scalability)
- Supporting a large range of devices and user interfaces
- Automated (provision-deploy-scale)
- CI/CD (agility)
- Support multiple datastore
- APIs at the heart of the applications
- Microservices



# Rules for developing/moving applications to the cloud

---

1. Don't code your application directly to a specific topology
2. Don't assume the local file system is permanent
3. Don't keep session state in your application
4. Don't log to the file system
5. Don't assume any specific infrastructure dependency
6. Don't use infrastructure APIs from within your application
7. Don't use obscure protocols
8. Don't rely on OS-specific features
9. Don't manually install your application

# Cloud-Native Application Goals – Day1

## Horizontal scaling

- Application runs in multiple runtimes spread across multiple hosts (VIII)

## Immutable deployment

- A runtime is not patched, it's replaced (IX)
- A runtime is stateless (VI)
- Shared functionality in backing services (IV)

## Elasticity

- Automatic scale-out and scale-in to maintain performance
- Achieved via containerization

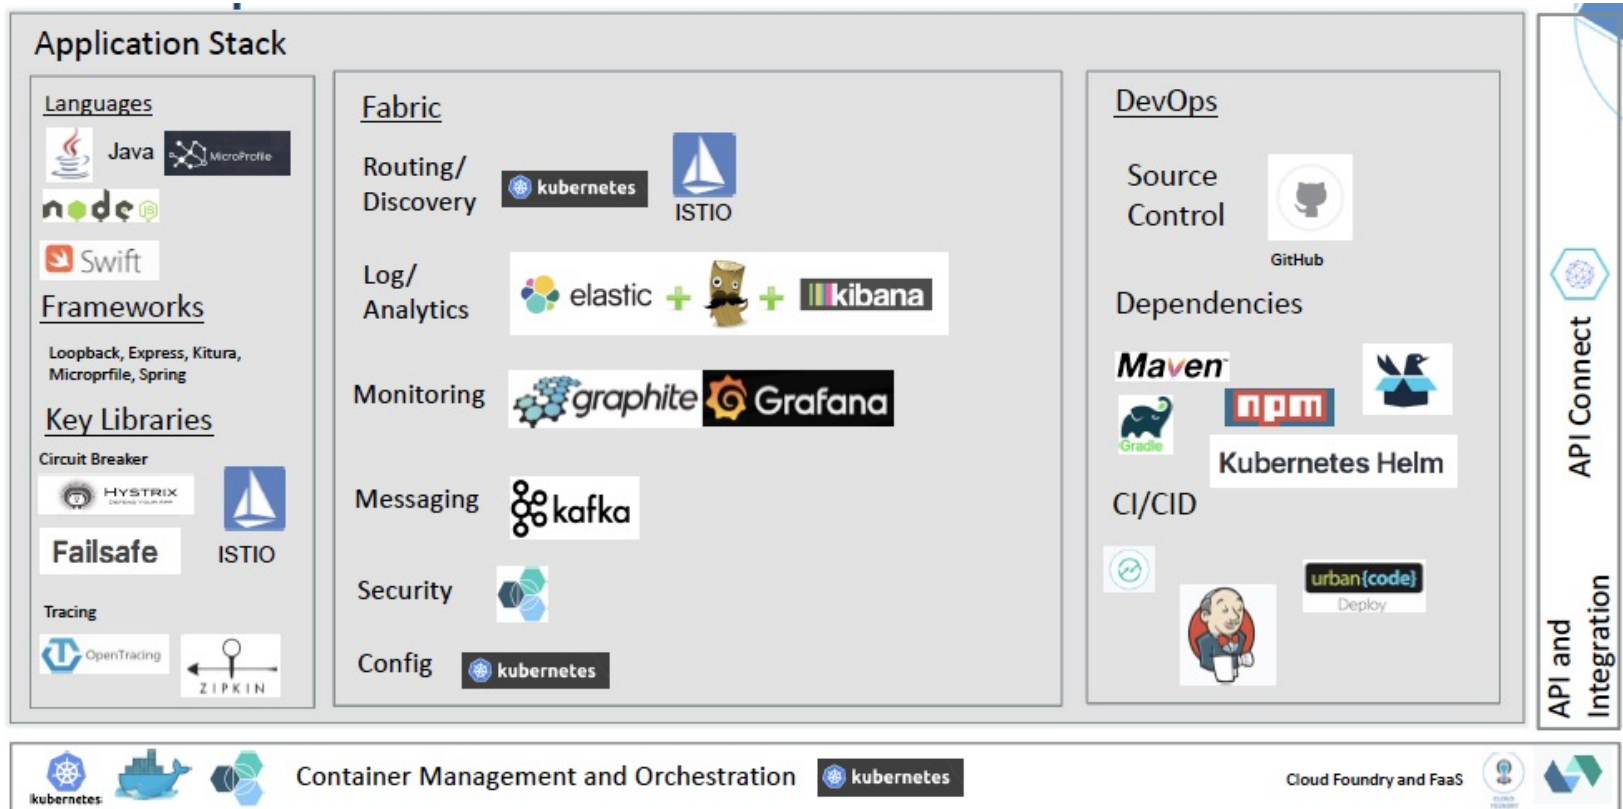
## Pay-as-you-go charging model

- Pay for what you use

### 12 factors for the Impatient

- I. Codebase - use version control (e.g. git)
- II. Dependencies - use a dependency manager (e.g. gradle/maven/sbt)
- III. Config - separate configuration from code (use the OS environment)
- IV. Backing Services - reference resources such as DBs by URLs in the config
- V. Build release run - separate build from run. Use versions.
- VI. Processes - run the app as one or more *stateless* processes.
- VII. Port binding - app should be self-contained. No app server.
- VIII. Concurrency - scale horizontally
- IX. Disposability - fast startup, graceful shutdown
- X. Dev/Prod parity - keep environments similar
- XI. Logs - treat logs as event streams (no FileAppenders!)
- XII. Admin Processes - treat admin processes as one-off events

# Development Stack - Choices



# Microservices: Making developers more efficient

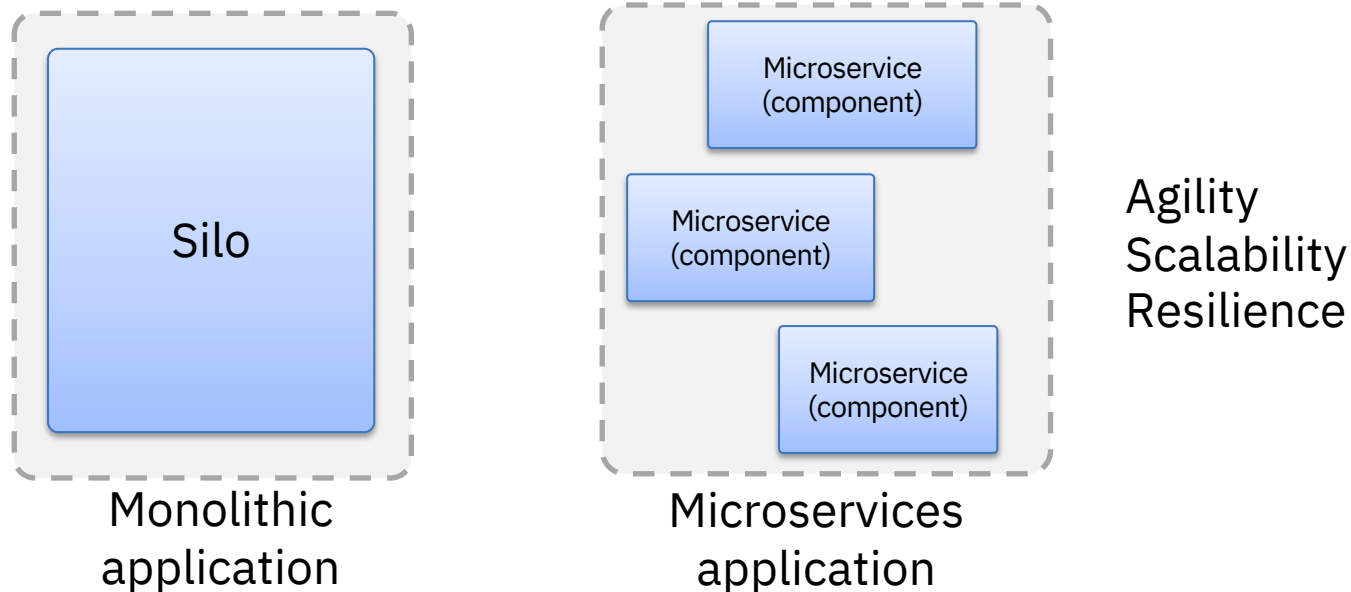
---

- An **engineering approach** that reduces an application into **single-function** modules
- They have well-defined **interfaces** that are independently deployed
- They are operated by a **small team** which owns the entire **lifecycle** of the service
- Microservices accelerate delivery by
  - minimizing communication and coordination between people
  - reducing the scope and risk of change



# Microservices Architecture ?

Simplistically, microservices architecture is about breaking down large silo applications into more manageable fully decoupled pieces



A microservice is a granular decoupled component within a broader application

# Why Microservices?

Small scoped, independent, scalable components

## Scaling

Elastic scalability

Workload orchestration

## Agility


Faster iteration cycles

Bounded context (code and data)

## Resilience

Reduced dependencies

Fail fast



# Microservices inter-communication

Aim is decoupling for robustness

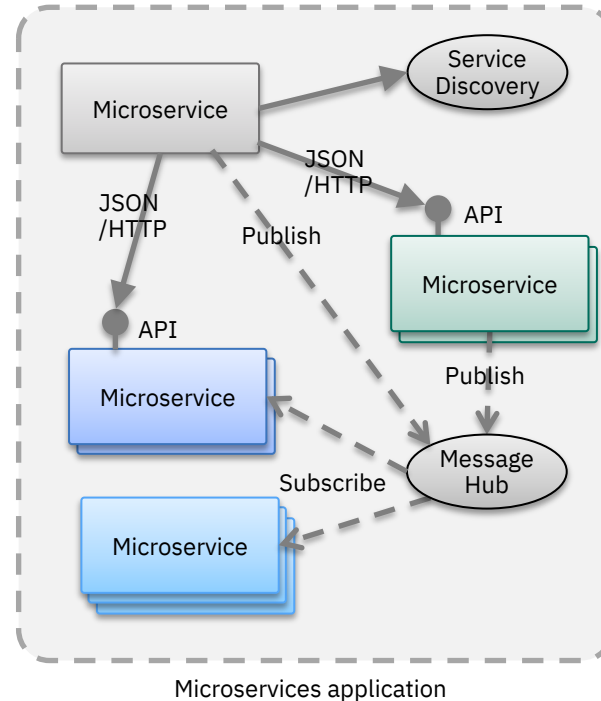
Compose a complex application  
using

“small”

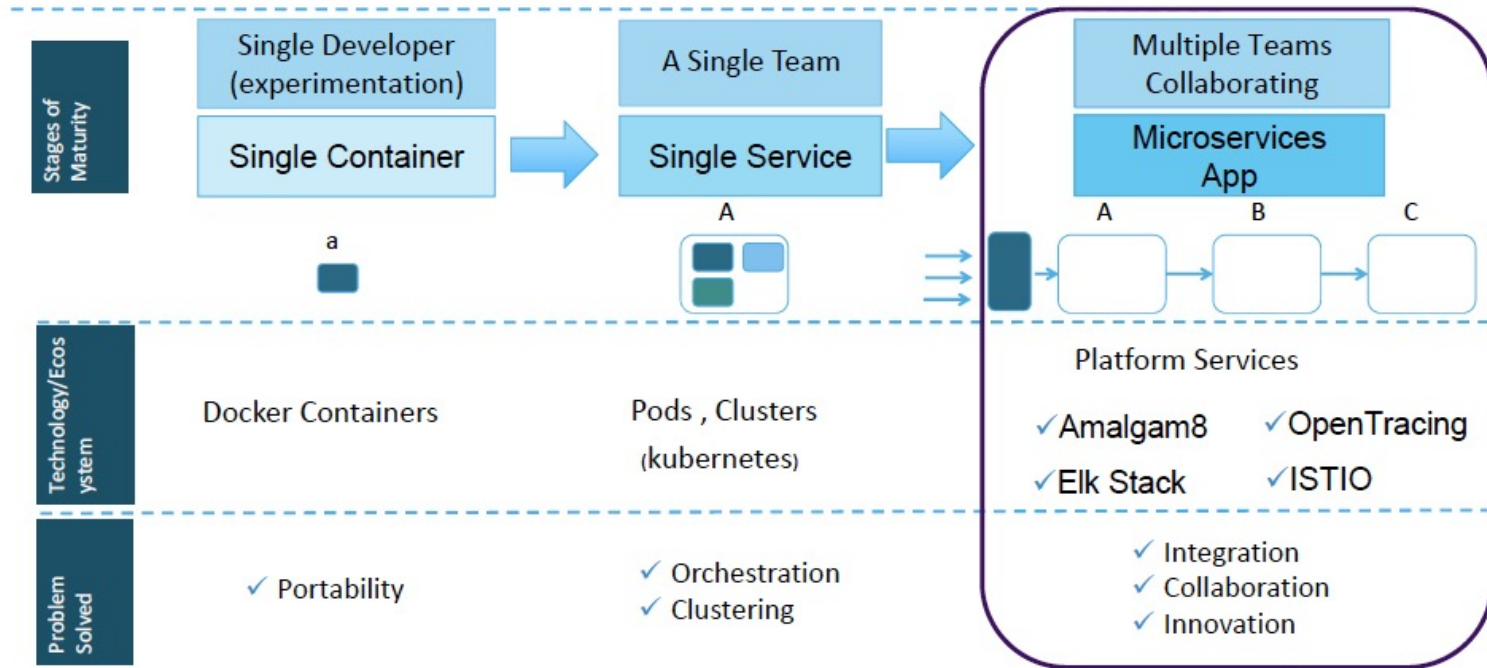
independent (autonomous)  
replaceable  
processes

that communicate

via language-agnostic APIs  
synchronously and asynchronously



# Microservices: Built for the enterprise journey



# Advantages / Challenges of Microservices

## Advantages

- Developed independently
- Developed by a single team
- Developed on its own timetable
- Each can be developed in a different language
- Manages its own data
- Scales and fails independently

## Challenges

- Developers must have significant operational skills (DevOps)
- Service interfaces and versions
- Duplication of effort across service implementations
- Extra complexity of creating a distributed system:

Designing decoupled, non-transactional systems is difficult

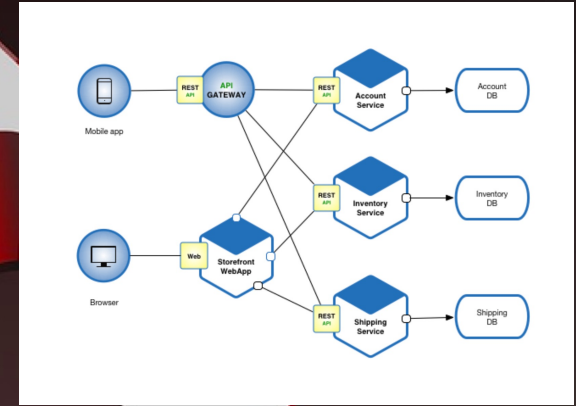
Locating service instances

Maintaining availability and consistency with partitioned data

End-to-end testing

# Things to take care of

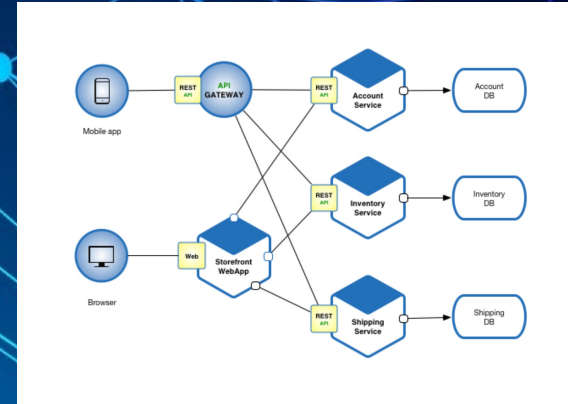
- Service discovery
- Load balancing
- Fault Tolerance
- Monitoring
- Dynamic routing
- New versions
- Security



# What is a 'Service Mesh' ?

A network for services

- Observability
- Resiliency
- Traffic Control
- Security
- Policy Enforcement
- Zero code change




# Istio, the service mesh – Day 2



## Istio

Connect, secure, control, and observe services.




Connect

Intelligently control the flow of traffic and API calls between services, conduct a range of tests, and upgrade gradually with red/black deployments.




Secure

Automatically secure your services through managed authentication, authorization, and encryption of communication between services.



Control

Apply policies and ensure that they're enforced, and that resources are fairly distributed among consumers.






Observe

See what's happening with rich automatic tracing, monitoring, and logging of all your services.

- IBM
- Lyft
- Google
- Some others

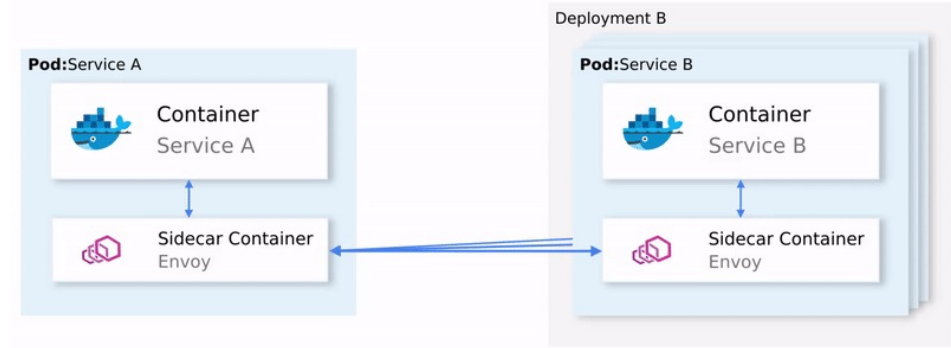


# What is Service Mesh?

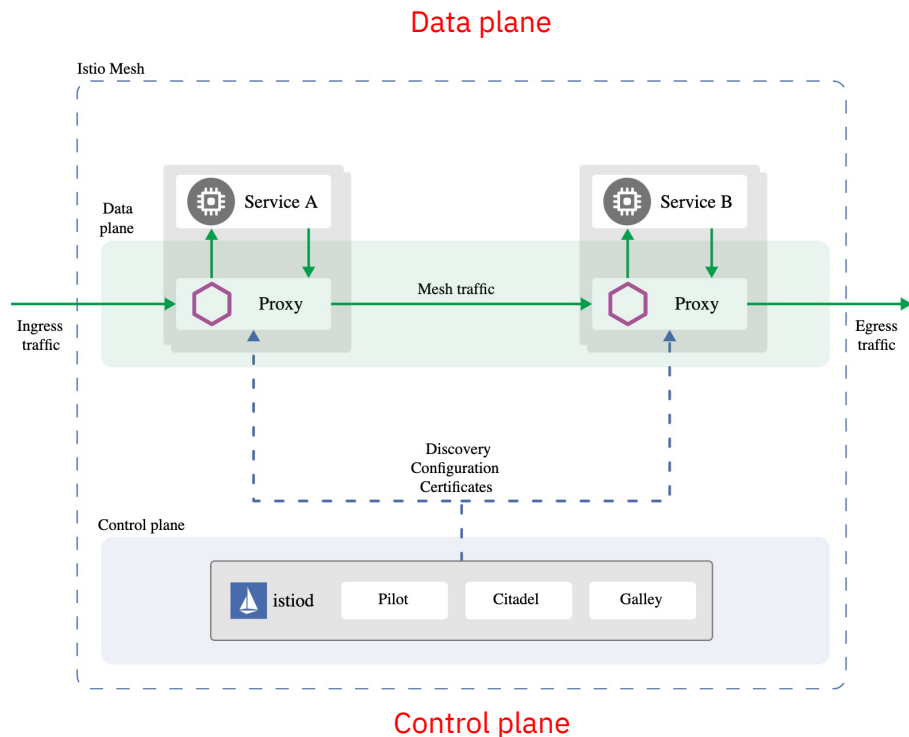


# Envoy : the sidecar proxy in each POD

- Dynamic service discovery
- Load balancing
- TLS termination
- HTTP/2 and gRPC proxies
- Circuit breakers
- Health checks
- Staged rollouts with %-based traffic split
- Fault injection
- Rich metrics




# Istio Architecture



- The **data plane** is composed of a set of intelligent proxies ([Envoy](#)) deployed as sidecars. These proxies mediate and control all network communication between microservices.
- The **control plane** manages and configures the proxies to route traffic.

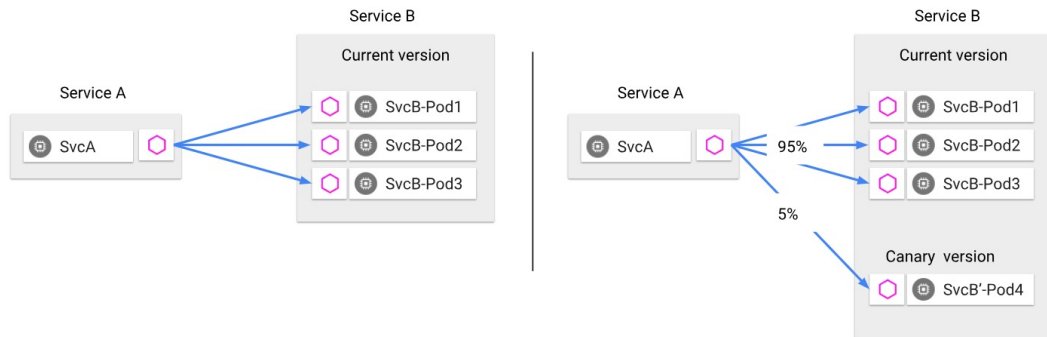
# Performance and Scalability (Istio 1.9.5)

- 1000 services, 2000 sidecars
- 70000 requests per second

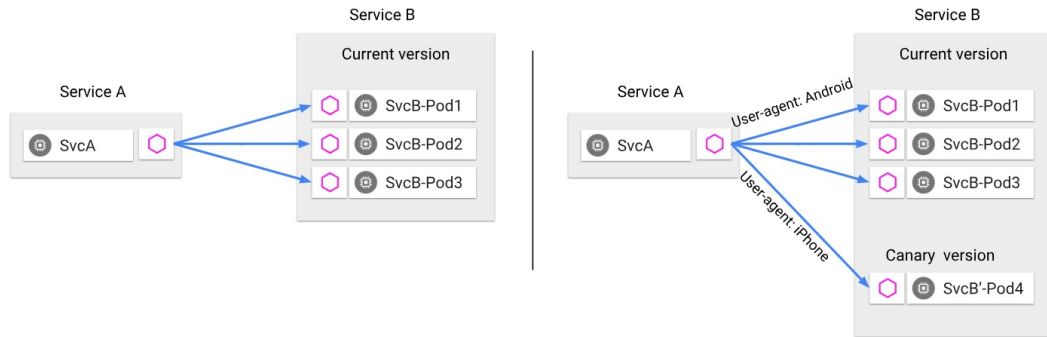


- **Envoy** proxy uses **0.35 vCPU** and **40 MB** memory per 1000 requests per second
- **Envoy** proxy adds **2.65 ms** to the 90th percentile latency.
- **Istiod** uses **1 vCPU** and 1.5 GB of memory.
- **Istiod** can be scaled horizontally

# Traffic Management



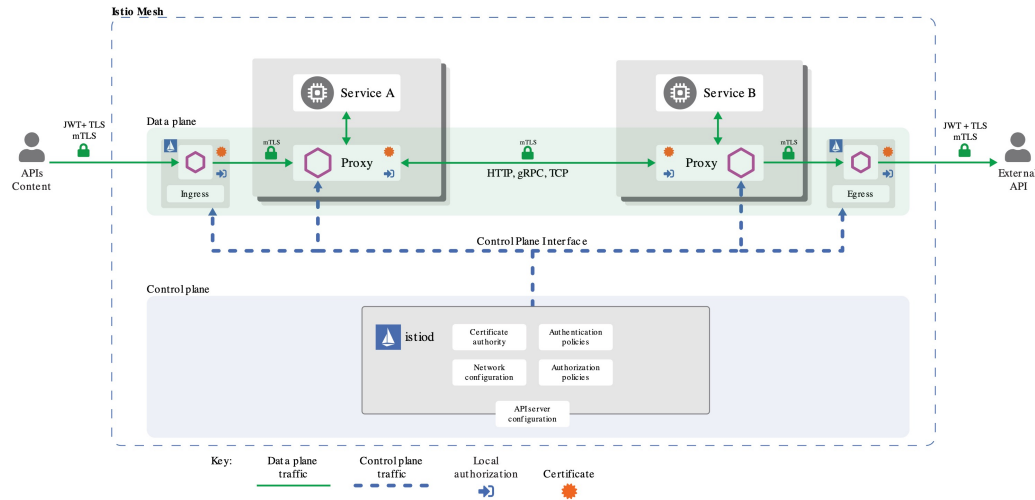
**Traffic splitting decoupled from infrastructure scaling** - proportion of traffic routed to a version is independent of number of instances supporting the version



**Content-based traffic steering** - The content of a request can be used to determine the destination of a request

- Request Routing
- Fault Injection
- Traffic Shifting
- TCP Traffic Shifting
- Request Timeouts
- Circuit Breaking
- Mirroring
- Rate limits

# Security

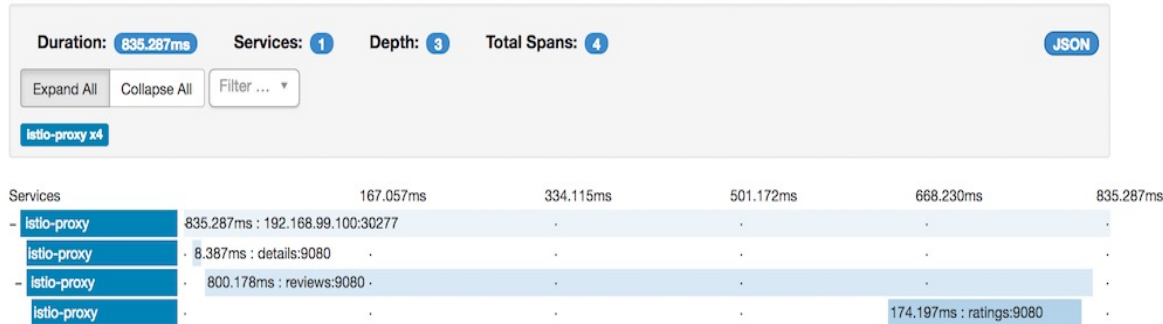


- Certificate Management
- Authentication
- Authorization

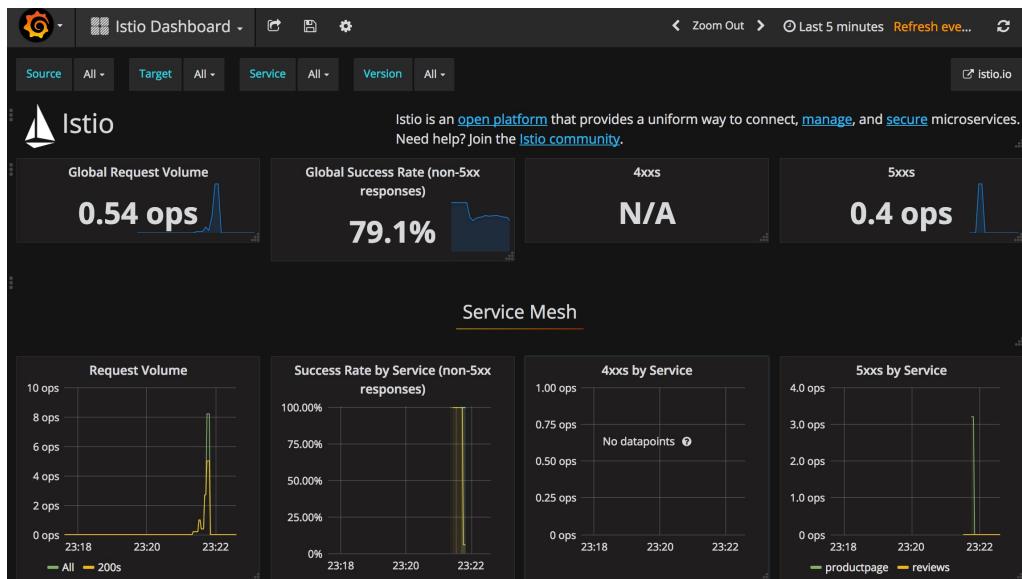
# Observability

Zipkin Investigate system behavior Find a trace Dependencies

Go to trace



- Metrics
- Logs
- Distributed Tracing



# Istio



**Istio runs in  
Production  
@WeatherCompan  
y**

## Business

- Weather forecasting

## Workload

- Manages `api.weather.com`
- 40 backend services
- 400K req/sec across the world
- 93% of reqs get...